

Digital Design & Computer Arch.

Lecture 12: Pipelined Processor Design II

Dr. Mohammad Sadrosadati

Prof. Onur Mutlu

ETH Zürich

Spring 2026

27 March 2026

Agenda for Today & Next Few Lectures

- Instruction Set Architectures (ISA): LC-3 and MIPS
- Assembly programming: LC-3 and MIPS
- Microarchitecture (principles & single-cycle uarch)
- Multi-cycle microarchitecture
- **Pipelining**
- **Issues in Pipelining:**
 - ❑ **Control & Data Dependence Handling**
 - ❑ State Maintenance and Recovery
- Out-of-Order Execution

Problem
Algorithm
Program/Language
System Software
SW/HW Interface
Micro-architecture
Logic
Devices
Electrons

Readings


■ This week

- Introduction to microarchitecture and single-cycle microarchitecture
 - H&H, Chapter 7.1-7.3
 - P&P, Appendices A and C
- Multi-cycle microarchitecture
 - H&H, Chapter 7.4
 - P&P, Appendices A and C


■ Also this week

- Pipelining
 - H&H, Chapter 7.5
- Pipelining Issues
 - H&H, Chapter 7.7, 7.8.1-7.8.3


Review: Single-Cycle MIPS Microarchitecture



Review: Multi-Cycle MIPS Processor



Review: Multi-Cycle MIPS FSM



**We can easily
handle multi-cycle
memory access**


Can We Do Better?

Pipelining

Example: Execution of Four Independent ADDs


- Multi-cycle: 4 cycles per instruction

1 instruction completed every 4 cycles




- Pipelined: 4 cycles per 4 instructions (steady state)

1 instruction completed every 1 cycle





Is life always this beautiful?

The Laundry Analogy




- “place one dirty load of clothes in the washer”
 - “when the washer is finished, place the wet load in the dryer”
 - “when the dryer is finished, take out the dry load and fold”
 - “when folding is finished, put the clothes away”
- steps to do a load are sequentially dependent
 - no dependence between different loads
 - different steps do not require the same resource(s)

Pipelining Multiple Loads of Laundry





- 4 loads of laundry in parallel (in steady state)
- **throughput** increased by 4x
- **latency** per load is the same
- no additional resources

Pipelining Multiple Loads of Laundry: In Practice



the slowest step (the dryer) decides throughput

Pipelining Multiple Loads of Laundry: In Practice



A Real-Life Pipeline: Automobile Assembly



From the collections of the Henry Ford

The guinea pig was the T's magneto, a component that supplied ignition energy to the engine before generators became common. A complex and innovative component that was one of the early Model T's technological advantages, Ford's magneto was integrated with the engine's flywheel and involved many pieces. Under the old system, each magneto was assembled by one worker. On average, that worker could assemble 35 of them in a nine-hour shift, or roughly one every 15 minutes.



Ford's transition to moving assembly lines began in April 1913 with the integrated (and complex) flywheel/magneto. With each worker assigned to complete a few specific tasks rather than build the entire unit, Ford reduced magneto assembly time from about 15 minutes to 5, and the required workforce decreased from 29 to 14.

A Real-Life Pipeline: Automobile Assembly



An Old Pipelined Computer: IBM Stretch



Design

Manufacturer	IBM
Designer	Gene Amdahl
Release date	May 1961
Units sold	9
Price	US\$7,780,000 (equivalent to \$67,380,000 in 2020)

Casing

Weight	70,000 pounds (35 short tons; 32 t) ^[1]
Power	100 kW ^[1] @ 110 V

System

Operating system	MCP
CPU	64-bit processor
Memory	2048 kilobytes (262144 x 64bits) ^[1]
MIPS	1.2 MIPS




V•T•E

An Ideal Pipeline

- Goal: **Increase throughput with little increase in cost** (hardware cost, in case of instruction processing)
- Repetition of **identical operations**
 - ❑ The same operation is repeated on a large number of different inputs (e.g., all laundry loads go through the same steps)
- Repetition of **independent operations**
 - ❑ No dependences between repeated operations
- **Uniformly partitionable suboperations**
 - ❑ Processing can be evenly divided into uniform-latency suboperations (that do not share resources)
- Fitting examples: automobile assembly line, doing laundry
 - ❑ What about the instruction processing “cycle”?

Ideal Pipelining

T_{put} = Throughput



More Realistic Pipeline: Throughput

- Nonpipelined version with delay T

$T_{\text{put}} = 1 / (T + S)$ where S = register (sequential logic) delay




- k-stage pipelined version

$$T_{\text{put}_{k\text{-stage}}} = 1 / (T/k + S)$$

$$T_{\text{put}_{\text{max}}} = 1 / (1 \text{ gate delay} + S)$$

**Register delay reduces throughput
(sequencing overhead b/w stages)**



This picture assumes “perfect division of work between stages (T/k)”

More Realistic Pipeline: Cost

- Nonpipelined version with combinational cost G


$\text{Cost} = G + R$ where R = register cost



- k -stage pipelined version

$\text{Cost}_{k\text{-stage}} = G + Rk$


Registers increase hardware cost




This picture ignores resource and register replication that is likely needed (G/k and Rk)

Pipelining Instruction Processing

Remember: The Instruction Processing Cycle




Remember: The Instruction Processing Cycle




1. Instruction fetch (IF)
2. Instruction decode and register operand fetch (ID/RF)
3. Execute/Evaluate memory address (EX/AG)
4. Memory operand fetch (MEM)
5. Store/writeback result (WB)

Remember the Single-Cycle Microarchitecture





Dividing the Single-Cycle Uarch Into Stages



Is this the correct partitioning?
Why not 4 or 6 stages? Why not different boundaries?



Instruction Pipeline Throughput



5-stage speedup is 4, not 5 as predicted by the ideal model. Why?

Half of the clock cycle is wasted in two stages (ID/RF and WB)


Enabling Pipelined Processing: Pipeline Registers




Pipelined Operation Example

All instruction classes must follow the same path and timing through the pipeline stages


Any performance impact?



Pipelined Operation Example




Illustrating Pipeline Operation: Operation View



Illustrating Pipeline Operation: Resource View

	t ₀	t ₁	t ₂	t ₃	t ₄	t ₅	t ₆	t ₇	t ₈	t ₉	t ₁₀
IF	I ₀	I ₁	I ₂	I ₃	I ₄	I ₅	I ₆	I ₇	I ₈	I ₉	I ₁₀
ID		I ₀	I ₁	I ₂	I ₃	I ₄	I ₅	I ₆	I ₇	I ₈	I ₉
EX			I ₀	I ₁	I ₂	I ₃	I ₄	I ₅	I ₆	I ₇	I ₈
MEM				I ₀	I ₁	I ₂	I ₃	I ₄	I ₅	I ₆	I ₇
WB					I ₀	I ₁	I ₂	I ₃	I ₄	I ₅	I ₆

Control Points in a Pipeline




Based on original figure from [P&H CO&D,
COPYRIGHT 2004 Elsevier. ALL RIGHTS
RESERVED.]

Identical set of control points as the single-cycle datapath

Control Signals in a Pipeline


- For a given instruction
 - ❑ same control signals as single-cycle, but
 - ❑ control signals required at different cycles, depending on stage
- ⇒ Option 1: decode once using the same logic as single-cycle and buffer signals until consumed




- ⇒ Option 2: carry relevant “instruction word/field” down the pipeline and decode locally within each or in a previous stage

Which one is better?


Pipelined Control Signals




Another Example: Single-Cycle and Pipelined



Is this correct?




Another Example: **Correct** Pipelined Datapath



- **WriteReg control signal must arrive at the same time as Result**

Another Example: Pipelined Control



- Same control unit as single-cycle processor
Control delayed to proper pipeline stage

Pipelined Datapath and Control Logic

■ Pipelined datapath design


- ❑ Start with the single-cycle datapath
- ❑ Decide how to break it down into stages
- ❑ Add pipeline registers to separate stages from each other
- ❑ Propagate data signals to stages where they are needed
- ❑ Make sure correct action is done in the correct stage
 - Wires that cross stages need to be handled carefully

■ Pipelined control logic design

- ❑ Start with the single-cycle control signals and control logic
- ❑ Propagate each control signal to stages where it is needed
 - Need to extend pipeline registers to do so

Is this sufficient?

Basic Pipelined Processor



How do we handle dependent instructions?

How do we handle multi-cycle memory access?

Remember: An Ideal Pipeline

- Goal: **Increase throughput with little increase in cost** (hardware cost, in case of instruction processing)
- Repetition of **identical operations**
 - ❑ The same operation is repeated on a large number of different inputs (e.g., all laundry loads go through the same steps)
- Repetition of **independent operations**
 - ❑ No dependencies between repeated operations
- **Uniformly partitionable suboperations**
 - ❑ Processing can be evenly divided into uniform-latency suboperations (that do not share resources)
- Fitting examples: automobile assembly line, doing laundry
 - ❑ **What about the instruction processing "cycle"?**

Instruction Pipeline: Not An Ideal Pipeline

- Identical operations ... NOT!

- ⇒ different instructions → not all need the same stages

- Forcing different instructions to go through the same pipe stages

- external fragmentation (some pipe stages idle for some instructions)

- Uniform suboperations ... NOT!

- ⇒ different pipeline stages → not the same latency

- Need to force each stage to be controlled by the same clock

- internal fragmentation (some pipe stages are fast but still have to take the same clock cycle time)

- Independent operations ... NOT!

- ⇒ instructions are not independent of each other

- Need to detect and resolve inter-instruction dependences to ensure the pipeline provides correct results

- pipeline stalls (pipeline is not always moving)

Issues in Pipeline Design

- Balancing work in pipeline stages
 - How many stages and what is done in each stage
- Keeping the pipeline correct, moving, and full in the presence of events that disrupt pipeline flow
 - Handling dependences
 - Data
 - Control
 - Handling resource contention
 - Handling long-latency (multi-cycle) operations
- Handling exceptions, interrupts
- Advanced: Improving pipeline throughput
 - Minimizing *stalls*

Causes of Pipeline *Stalls*

- Stall: A condition when the pipeline stops moving
- Resource contention
- Dependences (between instructions)
 - ❑ Data
 - ❑ Control
- Long-latency (multi-cycle) operations

Dependences and Their Types

- Also called “dependency” or *less desirably* “hazard”
- Dependences dictate ordering requirements between instructions
- Two types
 - Data dependence
 - Control dependence
- Resource contention is sometimes called resource dependence
 - However, this is not fundamental to (dictated by) program semantics, so we will treat it separately

Handling Resource Contention

- Happens when instructions in two pipeline stages need the same resource
- Solution 1: Eliminate the cause of contention
 - Duplicate the resource or increase its throughput
 - E.g., use separate instruction and data memories (caches)
 - E.g., use multiple ports for memory structures
- Solution 2: Detect the resource contention and stall one of the contending stages
 - Which stage do you stall?
 - Example: What if you had a single read and write port for the register file?

Example Resource Dependence: RegFile

- With careful design, the register file can be read and written in the same cycle:
 - write takes place during the 1st half of the cycle
 - read takes place during the 2nd half of the cycle => no problem!!!
 - read/write from/to the register file has only *half a clock cycle* to complete




Figure 7.44 Abstract view of pipeline in operation

Data Dependences

- Data dependence types
 - ❑ Flow dependence (true data dependence – read after write)
 - ❑ Anti dependence (write after read)
 - ❑ Output dependence (write after write)

- Which ones cause stalls in a pipelined machine?
 - ❑ For all of them, we need to ensure semantics of the program is correct
 - ❑ Flow dependences always need to be obeyed because they constitute true dependence on a value
 - ❑ Anti and output dependences exist due to limited number of architectural registers
 - They are dependence on a name, not a value
 - We will later see what we can do about them

Data Dependence Types

Flow dependence

$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_5 \leftarrow r_3 \text{ op } r_4$



Read-after-Write
(RAW)

Anti dependence

$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_1 \leftarrow r_4 \text{ op } r_5$



Write-after-Read
(WAR)


Output dependence

$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_5 \leftarrow r_3 \text{ op } r_4$
 $r_3 \leftarrow r_6 \text{ op } r_7$



Write-after-Write
(WAW)

Pipelined Operation Example



Data Dependence Handling

Readings for Next Few Lectures

- H&H, Chapter 7.5-7.9
- Smith and Sohi, “[The Microarchitecture of Superscalar Processors](#),” Proceedings of the IEEE, 1995
 - ❑ More advanced pipelining
 - ❑ Interrupt and exception handling
 - ❑ Out-of-order and superscalar execution concepts


How to Handle Data Dependences

- Anti and output dependences are easier to handle
 - ❑ write to the destination only in last stage and in program order
- Flow dependences are more interesting & challenging
- Six fundamental ways of handling flow dependences
 - ❑ Detect and wait until value is available in register file
 - ❑ Detect and forward/bypass data to dependent instruction
 - ❑ Detect and eliminate the dependence at the software level
 - No need for the hardware to detect dependence
 - ❑ Detect and move it out of the way for independent instructions
 - ❑ Predict the needed value(s), execute “speculatively”, and verify
 - ❑ Do something else (fine-grained multithreading)
 - No need to detect

Recall: Data Dependence Types

Flow dependence


$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_5 \leftarrow r_3 \text{ op } r_4$



Read-after-Write
(RAW)

Anti dependence


$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_1 \leftarrow r_4 \text{ op } r_5$



Write-after-Read
(WAR)

Output dependence


$r_3 \leftarrow r_1 \text{ op } r_2$
 $r_5 \leftarrow r_3 \text{ op } r_4$
 $r_3 \leftarrow r_6 \text{ op } r_7$




Write-after-Write
(WAW)

RAW Dependence Handling

- Which one of the following flow dependences lead to conflicts in the 5-stage pipeline?



Pipeline Stall: Resolving Data Dependence



Interlocking

- Interlocking: Detection of dependence between instructions in a pipelined processor to guarantee correct execution
- Software based interlocking
vs.
- Hardware based interlocking
- MIPS acronym?

Approaches to Dependence Detection (I)

■ Scoreboarding

- ❑ Each register in register file has a Valid bit associated with it
- ❑ An instruction that is writing to the register resets the Valid bit
- ❑ An instruction in Decode stage checks if all its source and destination registers are Valid
 - Yes: No need to stall... No dependence
 - No: Stall the instruction


■ Advantage:

- ❑ Simple. 1 bit per register

■ Disadvantage:

- ❑ Need to stall for all types of dependences, not only flow dep.

Pipelined Operation Example




Now assume SUB is dependent on LW

Approaches to Dependence Detection (II)

- **Combinational dependence check logic**
 - ❑ Special logic checks if any instruction in later stages is supposed to write to any source register of the instruction that is being decoded
 - ❑ Yes: stall the instruction/pipeline
 - ❑ No: no need to stall... no flow dependence
- **Advantage:**
 - ❑ No need to stall on anti and output dependences
- **Disadvantage:**
 - ❑ Logic is more complex than a scoreboard
 - ❑ Logic becomes more complex as we make the pipeline deeper and wider (flash-forward: superscalar execution, later lecture)

Pipelined Operation Example



Now assume SUB is dependent on LW

Once You Detect the Dependence in Hardware


- What do you do afterwards?
- Observation: Dependence between two instructions is detected before the communicated data value becomes available
- Option 1: Stall the dependent instruction right away
- Option 2: Stall the dependent instruction only when necessary → data forwarding/bypassing
- Option 3: ...

Data Forwarding/Bypassing

- Problem: A consumer (dependent) instruction has to wait in decode stage until the producer instruction writes its value in the register file
- Goal: We do not want to stall the pipeline unnecessarily
- Observation: The data value needed by the consumer instruction can be supplied directly from a later stage in the pipeline (instead of only from the register file)
- Idea: Add additional dependence check logic and data forwarding paths (buses) to supply the producer's value to the consumer right after the value is available
- Benefit: Consumer can move in the pipeline until the point the value can be supplied → **less stalling**

Data Dependence Handling: Concepts and Implementation

How to Implement Stalling



■ Stall


- ❑ disable **PC** and **IF/ID** latching; ensure stalled instruction stays in its stage
- ❑ Insert **"invalid"** instructions/nops into the stage following the stalled one (called **"bubbles"**)

RAW Data Dependence Example


- One instruction writes a register (\$s0) and next instructions read this register => read after write (RAW) dependence.

**Wrong results happen only if
the pipeline handles
data dependences incorrectly!**

- subsequent instructions read the correct value of \$s0



Compile-Time Detection and Elimination




- Insert enough independent instructions for the required result to be ready by the time it is needed by a dependent one
 - Reorder/reschedule/insert instructions at the compiler level

Data Forwarding

- Also called Data Bypassing
 - Forward the result value to the dependent instruction as soon as the value is available
 - We have already seen the basic idea before
 - Remember dataflow?
 - Data value is supplied to dependent instruction as soon as it is available
 - Instruction executes when all its operands are available
 - Data forwarding brings a pipeline closer to dataflow execution principles
-

Data Forwarding: Locations in Datapath




From latched output of ALU to input of ALU
From WB to input of ALU
From WB to RF (internal in Register File)

Data Forwarding: Datapath & Control

From latched output of ALU to input of ALU

From WB to input of ALU

From Regfile



Data Forwarding: Implementation

- Forward to Execute stage from either:
 - Memory stage or
 - Writeback stage
 - When should we forward from either Memory or Writeback stage?
 - If that stage will write to a destination register and the destination register matches the source register
 - If both the Memory & Writeback stages contain matching destination registers, Memory stage has priority to forward its data, because it contains the *more recently executed* instruction
-

Data Forwarding (in Pseudocode)

- Forward to Execute stage from either:


- Memory stage or
- Writeback stage

- Forwarding logic for *ForwardAE* (pseudo code):

```
if ((rsE != 0) AND (rsE == WriteRegM) AND RegWriteM) then
    ForwardAE = 10 # forward from Memory stage
else if ((rsE != 0) AND (rsE == WriteRegW) AND RegWriteW) then
    ForwardAE = 01 # forward from Writeback stage
else
    ForwardAE = 00 # no forwarding
```


- Forwarding logic for *ForwardBE* same, but replace *rsE* with *rtE*
-

Data Forwarding Is **Not** Always Possible




- Forwarding is usually sufficient to resolve RAW data dependences
- Unfortunately, there are cases when forwarding is **not possible**
 - due to pipeline design and instruction latencies
 - The `lw` instruction **does not finish** reading data until the end of Memory stage
→ its result **cannot be forwarded** to the Execute stage of the next instruction
unless we want a long critical path → **breaks critical path design principle**

Stalling Necessary for MEM-EX Dependence



Stalling and Dependence Detection Hardware



Hardware Needed for Stalling

- Stalls are supported by adding
 - enable inputs (EN) to the Fetch and Decode pipeline registers
 - synchronous reset/clear (CLR) input to the Execute pipeline register
 - or an INV bit associated with each pipeline register, indicating that contents are INValid

 - When a lw stall occurs
 - Keep the values in the Decode and Fetch stage pipeline registers
 - StallD and StallF are asserted
 - Clear the contents of the Execute stage register, introducing a bubble
 - FlushE is also asserted
-

A Special Case of Data Dependence

- Control dependence
 - Data dependence on the Instruction Pointer / Program Counter


Control Dependence

- Question: **What should the fetch PC be in the next cycle?**
- Answer: The address of the next instruction
 - ❑ All instructions are control dependent on previous ones. Why?
- If the fetched instruction is a non-control-flow instruction:
 - ❑ Next Fetch PC is the address of the next-sequential instruction
 - ❑ Easy to determine if we know the size of the fetched instruction
- **If the instruction that is fetched is a control-flow instruction:**
 - ❑ How do we determine the next Fetch PC?
- In fact, how do we know whether or not the fetched instruction is a control-flow instruction?


Branch Prediction

- Special case of data dependence: **dependence on PC**
- **beq:**
 - Conditional branch is not resolved until the fourth stage of the pipeline
 - **Instructions after the branch are fetched before branch is resolved**
 - Simple “**branch prediction**” example:
 - Always predict that the next sequential instruction is fetched
 - Called “**Always not taken**” prediction
 - Flush (invalidate) “not-taken path” instructions if the branch is taken
- **Branch misprediction penalty**
 - **number of instructions flushed when branch is incorrectly predicted**
 - Penalty can be reduced by resolving the branch earlier
 - Called “**Early branch resolution**”


Our Pipeline So Far



Control Dependence: Flush on Misprediction




Pipeline with Early Branch Resolution



Need to calculate branch target and condition in the Decode Stage

Early Branch Resolution



Early Branch Resolution: Good Idea?

■ Advantages

- Reduced branch misprediction penalty
 - Reduced CPI (cycles per instruction)


■ Disadvantages

- Potential increase in clock cycle time?
 - Higher clock period and lower frequency?
- Additional hardware cost
 - Specialized and likely not used by other instructions

Recall: Performance Analysis Basics

- Execution time of a single instruction
 - **{CPI} x {clock cycle time}**
 - CPI: Number of cycles it takes to execute an instruction
- Execution time of an entire program
 - Sum over all instructions [**{CPI} x {clock cycle time}**]
 - **{# of instructions} x {Average CPI} x {clock cycle time}**

Data Forwarding for Early Branch Resolution



Data forwarding for early branch resolution adds even more complexity

Forwarding and Stalling Hardware Control

```
// Forwarding logic:
assign ForwardAD = (rsD != 0) & (rsD == WriteRegM) & RegWriteM;
assign ForwardBD = (rtD != 0) & (rtD == WriteRegM) & RegWriteM;

//Stalling logic:
assign lwstall = ((rsD == rtE) | (rtD == rtE)) & MemtoRegE;

assign branchstall = (BranchD & RegWriteE &
                     (WriteRegE == rsD | WriteRegE == rtD))
                     |
                     (BranchD & MemtoRegM &
                     (WriteRegM == rsD | WriteRegM == rtD));

// Stall signals;
assign StallF = lwstall | branchstall;
assign StallD = lwstall | branchstall;
assign FlushE = lwstall | branchstall;
```

Final Pipelined MIPS Processor (H&H)





Figure 7.58 Pipelined processor with full hazard handling

Includes always-not-taken br prediction, early branch resolution, forwarding, stall logic

Doing Better: Smarter Branch Prediction

- **Guess whether or not branch will be taken**
 - Backward branches are usually taken (loops iterate many times)
 - History of whether branch was previously taken can improve the guess
- **Accurate branch prediction reduces the fraction of branches requiring a flush**
- **Many sophisticated techniques are employed in modern processors**
 - Including simple machine learning methods (perceptrons)
 - We will see them in Branch Prediction lectures

More on Branch Prediction (I)



The video player displays a lecture slide titled "Importance of The Branch Problem". The slide contains the following text:

- Assume $N = 20$ (20 pipe stages), $W = 5$ (5 wide fetch)
- Assume: 1 out of 5 instructions is a branch
- Assume: Each 5 instruction-block ends with a branch
- How long does it take to fetch 500 instructions?
 - 100% accuracy
 - 100 cycles (all instructions fetched on the correct path)
 - No wasted work; $IPC = 500/100$
 - 99% accuracy
 - $100 \text{ (correct path)} + 20 * 1 \text{ (wrong path)} = 120 \text{ cycles}$
 - 20% extra instructions fetched; $IPC = 500/120$
 - 90% accuracy
 - $100 \text{ (correct path)} + 20 * 10 \text{ (wrong path)} = 300 \text{ cycles}$
 - 200% extra instructions fetched; $IPC = 500/300$
 - 60% accuracy
 - $100 \text{ (correct path)} + 20 * 40 \text{ (wrong path)} = 900 \text{ cycles}$
 - 800% extra instructions fetched; $IPC = 500/900$

The video player interface shows the video is at 19:43 / 1:05:13. Below the video, the title "Digital Design & Computer Architecture - Lecture 16b: Branch Prediction I (ETH Zürich, Spring 2020)" is displayed, along with 2,612 views and the date Apr 20, 2020. The channel "Onur Mutlu Lectures" is listed with 15.5K subscribers. The video description includes the course name, date, and a link to the lecture page.

ETH ZÜRICH
Digital Design & Computer Architecture - Lecture 16b: Branch Prediction I (ETH Zürich, Spring 2020)
2,612 views • Apr 20, 2020

Onur Mutlu Lectures
15.5K subscribers

Digital Design and Computer Architecture, ETH Zürich, Spring 2020
(<https://safari.ethz.ch/digitaltechnik...>)

Lecture 16b: Branch Prediction I
Lecturer: Professor Onur Mutlu (<https://people.inf.ethz.ch/omutlu/>)
Date: April 23, 2020

47 0 SHARE SAVE ...
ANALYTICS EDIT VIDEO

https://www.youtube.com/watch?v=h6l9yYSyZHM&list=PL5Q2soXY2Zi_FRrIoMa2fUYWPGiZUBQo2&index=22

More on Branch Prediction (II)

Intel Pentium M Predictors: Loop and Jump

The advanced branch prediction in the Pentium M processor is based on the Intel Pentium® 4 processor's [6] branch predictor. On top of that, two additional predictors to capture special program flows, were added: a Loop Detector and an Indirect Branch Predictor.




Figure 2: The Loop Detector logic




Figure 3: The Indirect Branch Predictor logic

Gochman et al.,
"The Intel Pentium M Processor: Microarchitecture and Performance,"
Intel Technology Journal, May 2003.

39

ETH ZÜRICH

Digital Design & Computer Architecture - Lecture 17: Branch Prediction II (ETH Zürich, Spring 2020)

2,846 views • Apr 20, 2020

42 0 SHARE SAVE ...



Onur Mutlu Lectures
15.5K subscribers

Digital Design and Computer Architecture, ETH Zürich, Spring 2020
(<https://safari.ethz.ch/digitaltechnik...>)


Lecture 17: Branch Prediction II
Lecturer: Professor Onur Mutlu (<https://people.inf.ethz.ch/omutlu/>)
Date: April 24, 2020

ANALYTICS

EDIT VIDEO

https://www.youtube.com/watch?v=z77VpggShvg&list=PL5Q2soXY2Zi_FRloMa2fUYWPGiZUBQo2&index=23

More on Branch Prediction (III)



18-740 Computer Architecture - Advanced Branch Prediction - Lecture 5

4,696 views • Sep 23, 2015

41 0 SHARE SAVE ...



Carnegie Mellon Computer Architecture
23K subscribers

Lecture 5: Advanced Branch Prediction

Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)

Date: September 16, 2014.

Lecture 5 slides (pdf): <http://www.ece.cmu.edu/~ece740/f15/li...>

Lecture 5 slides (ppt): <http://www.ece.cmu.edu/~ece740/f15/li...>

SUBSCRIBED



<https://www.youtube.com/watch?v=yDjsr-jT0tk&list=PL5PHm2jkkXmgVhh8CHAu9N76TShJqfYDt&index=4>

Lectures on Branch Prediction

- **Digital Design & Computer Architecture, Spring 2020, Lecture 16b**
 - Branch Prediction I (ETH Zurich, Spring 2020)
 - https://www.youtube.com/watch?v=h6l9yYSyZHM&list=PL5Q2soXY2Zi_FRrloMa2fUYWPGiZUBQo2&index=22

- **Digital Design & Computer Architecture, Spring 2020, Lecture 17**
 - Branch Prediction II (ETH Zurich, Spring 2020)
 - https://www.youtube.com/watch?v=z77VpggShvg&list=PL5Q2soXY2Zi_FRrloMa2fUYWPGiZUBQo2&index=23

- **Computer Architecture, Spring 2015, Lecture 5**
 - Advanced Branch Prediction (CMU, Spring 2015)
 - <https://www.youtube.com/watch?v=yDjsr-jTOtk&list=PL5PHm2jkkXmgVhh8CHAU9N76TShJqfYDt&index=4>

We Covered Until Here in Lecture

Digital Design & Computer Arch.

Lecture 12: Pipelined Processor Design II

Dr. Mohammad Sadrosadati

Prof. Onur Mutlu

ETH Zürich

Spring 2026

27 March 2026

**Further Slides for Your Own Study
(May Be Covered in Future Lectures)**

Pipelined Performance Example

Pipelined Performance Example

- **An important program consists of:**

- 25% loads
- 10% stores
- 11% branches
- 2% jumps
- 52% R-type

- **Assume:**

- 40% of loads used by next instruction
- 25% of branches mispredicted; flushes the next instruction on misprediction
- All jumps flush the next instruction fetched

- **What is the average CPI?**

Pipelined Performance Example: CPI

- Load/Branch CPI = 1 when no stall/flush, 2 when stall/flush.

Thus:

- $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$

Average CPI for load

- $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$

Average CPI for branch

- And

- *Average CPI* =

Pipelined Performance Example: CPI

- Load/Branch CPI = 1 when no stall/flush, 2 when stall/flush.

Thus:

- $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$

Average CPI for load

- $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$

Average CPI for branch

- And

- *Average CPI* = $(0.25)(1.4) +$
 $(0.1)(1) +$
 $(0.11)(1.25) +$
 $(0.02)(2) +$
 $(0.52)(1)$

load
store
beq
jump
r-type

= **1.15**

Pipelined Performance Example: Cycle Time

- There are 5 stages, and 5 different timing paths:

$$T_c = \max \{$$

$$t_{pcq} + t_{mem} + t_{setup}$$

$$2(t_{Rfread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup})$$

$$t_{pcq} + t_{mux} + t_{mux} + t_{ALU} + t_{setup}$$

$$t_{pcq} + t_{memwrite} + t_{setup}$$

$$2(t_{pcq} + t_{mux} + t_{RFwrite})$$

$$\}$$

fetch

decode

execute

memory

writeback

- The clock cycle *depends* on the *slowest stage*
- Decode and Writeback use register file and have only half a clock cycle to complete → that is why there is a 2 in front of them

Final Pipelined MIPS Processor (H&H)




Figure 7.58 Pipelined processor with full hazard handling

Includes always-not-taken br prediction, early branch resolution, forwarding, stall logic

Pipelined Performance Example: Cycle Time

Element	Parameter	Delay (ps)
Register clock-to-Q	t_{pcq_PC}	30
Register setup	t_{setup}	20
Multiplexer	t_{mux}	25
ALU	t_{ALU}	200
Memory read	t_{mem}	250
Register file read	t_{RFread}	150
Register file setup	$t_{RFsetup}$	20
Equality comparator	t_{eq}	40
AND gate	t_{AND}	15
Memory write	$T_{memwrite}$	220
Register file write	$t_{RFwrite}$	100

$$\begin{aligned}T_c &= 2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup}) \\&= 2[150 + 25 + 40 + 15 + 25 + 20] \text{ ps} \\&= 550 \text{ ps}\end{aligned}$$

Pipelined Performance Example: Exec Time

- For a program with 100 billion instructions executing on a pipelined MIPS processor:

- $CPI = 1.15$
- $T_c = 550 \text{ ps}$

- Execution Time $= (\# \text{ instructions}) \times CPI \times T_c$
 $= (100 \times 10^9)(1.15)(550 \times 10^{-12})$
 $= 63 \text{ seconds}$

Performance Summary for 3 MIPS microarch.

Processor	Execution Time (seconds)	Normalized Execution Time (single-cycle is baseline)
Single-cycle	95	1
Multicycle	133	0.71
Pipelined	63	1.51

- **Pipelined implementation is the fastest of 3 implementations**
- **Even though we have a 5-stage pipeline, speedup is not 5X over the single-cycle or the multi-cycle system!**

Recall: How to Handle Data Dependences

- Anti and output dependences are easier to handle
 - ❑ write to the destination only in last stage and in program order
- Flow dependences are more interesting
- Six fundamental ways of handling flow dependences
 - ❑ Detect and wait until value is available in register file
 - ❑ Detect and forward/bypass data to dependent instruction
 - ❑ Detect and eliminate the dependence at the software level
 - No need for the hardware to detect dependence
 - ❑ Detect and move it out of the way for independent instructions
 - ❑ Predict the needed value(s), execute “speculatively”, and verify
 - ❑ Do something else (fine-grained multithreading)
 - No need to detect

Recall: How to Handle Data Dependences

- Anti and output dependences are easier to handle
 - ❑ write to the destination only in last stage and in program order
- Flow dependences are more interesting
- Six fundamental ways of handling flow dependences
 - ❑ Detect and wait until value is available in register file
 - ❑ Detect and forward/bypass data to dependent instruction
 - ❑ Detect and eliminate the dependence at the software level
 - No need for the hardware to detect dependence
 - ❑ Detect and move it out of the way for independent instructions
 - ❑ Predict the needed value(s), execute “speculatively”, and verify
 - ❑ Do something else (fine-grained multithreading)
 - No need to detect

Question to Ponder: Hardware vs. Software

- What is the role of the hardware vs. the software in data dependence handling?
 - ❑ Software based interlocking
 - ❑ Hardware based interlocking
 - ❑ Who inserts/manages the pipeline bubbles?
 - ❑ Who finds the independent instructions to fill “empty” pipeline slots?
 - ❑ What are the advantages/disadvantages of each?
 - Think of the performance equation as well

Question to Ponder: Hardware vs. Software

- What is the role of the hardware vs. the software in the order in which instructions are executed in the pipeline?
 - ❑ Software based instruction scheduling → static scheduling
 - ❑ Hardware based instruction scheduling → dynamic scheduling
- How does each impact different metrics?
 - ❑ Performance (and parts of the performance equation)
 - ❑ Complexity
 - ❑ Power consumption
 - ❑ Reliability
 - ❑ Cost
 - ❑ ...

More on Software vs. Hardware

- Software based scheduling of instructions → static scheduling
 - ❑ Hardware executes the **instructions in the compiler-dictated order**
 - ❑ Contrast this with **dynamic scheduling**: hardware can execute instructions out of the compiler-specified order
 - ❑ How does the compiler know the latency of each instruction?
- What information does the compiler not know that makes static scheduling difficult?
 - ❑ Answer: Anything that is determined at run time
 - **Variable-length operation latency, memory address, branch direction**
- How can the compiler alleviate this (i.e., estimate the unknown)?
 - ❑ Answer: Profiling (done statically or dynamically)

More on Static Instruction Scheduling

Trace Scheduling Idea

TRACE SCHEDULING LOOP-FREE CODE

43

58:18 / 1:41:17

Lecture 16. Static Instruction Scheduling - Carnegie Mellon - Comp. Arch. 2015 - Onur Mutlu

7,136 views • Feb 26, 2015

46 0 SHARE SAVE ...

Carnegie Mellon Computer Architecture
23K subscribers

Lecture 16: Static Instruction Scheduling
Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)
Date: Feb 23rd, 2015

Lecture 16 slides (pdf): <http://www.ece.cmu.edu/~ece447/s15/li...>

SUBSCRIBED


<https://www.youtube.com/watch?v=isBEVkljgGA&list=PL5PHm2jkkXmi5CxxI7b3JCL1TWybtDtKq&index=18>

Lectures on Static Instruction Scheduling

- **Computer Architecture, Spring 2015, Lecture 16**
 - Static Instruction Scheduling (CMU, Spring 2015)
 - <https://www.youtube.com/watch?v=isBEVkiJgGA&list=PL5PHm2jkkXmi5CxxI7b3JCL1TWybTDtKq&index=18>
- **Computer Architecture, Spring 2013, Lecture 21**
 - Static Instruction Scheduling (CMU, Spring 2013)
 - <https://www.youtube.com/watch?v=XdDUn2WtkRg&list=PL5PHm2jkkXmidJOd59REog9jDnPDTG6IJ&index=21>


Recall: Semantic Gap

- How close instructions & data types & addressing modes are to high-level language (HLL)



HW
Control
Signals

Easier mapping of HLL to ISA
Less work for software designer
More work for hardware designer
Optimization burden on HW




HW
Control
Signals

Harder mapping of HLL to ISA
More work for software designer
Less work for hardware designer
Optimization burden on SW

Recall: How to Change the Semantic Gap Tradeoffs

- Translate from one ISA into a different “implementation” ISA



An Example: Rosetta 2 Binary Translator

Rosetta 2 [\[edit\]](#)


In 2020, Apple announced Rosetta 2 would be bundled with [macOS Big Sur](#), to aid in the [Mac transition to Apple silicon](#). The software permits many applications compiled exclusively for execution on x86-64-based processors to be translated for execution on Apple silicon.^{[2][8]}

In addition to the [just-in-time](#) (JIT) translation support, Rosetta 2 offers [ahead-of-time compilation](#) (AOT), with the x86-64 code fully translated, just once, when an application without a universal binary is installed on an Apple silicon Mac.^[9]


Rosetta 2's performance has been praised greatly.^{[10][11]} In some benchmarks, x86-64-only programs performed better under Rosetta 2 on a Mac with an Apple M1 SOC than natively on a Mac with an Intel x86-64 processor. One of the key reasons why Rosetta 2 provides such high level of translation efficiency is the support of x86-64 [memory ordering](#) in Apple M1 SOC.^[12]

Although Rosetta 2 works for most software, some software doesn't work at all^[13] or is reported to be "sluggish".^[14] A lot of software can be made compatible with the new Macs by the vendor recompiling the software, often a simple task; while for some software (such as software that includes [assembly language](#) code, or that generates [machine code](#)), the changes to make them work aren't simple and cannot be automated.

Similar to the first version, Rosetta 2 does not normally require user intervention. When a user attempts to launch an x86-64-only application for the first time, macOS prompts them to install Rosetta 2 if it is not already available. Subsequent launches of x86-64 programs will execute via translation automatically. An option also exists to force a universal binary to run as x86-64 code through Rosetta 2, even on an ARM-based machine.^[15]



An Example: Rosetta 2 Binary Translator





Apple M1,
2021

Another Example: NVIDIA Denver

The Secret of Denver: Binary Translation & Code Optimization

As we alluded to earlier, NVIDIA's decision to forgo a traditional out-of-order design for Denver means that much of Denver's potential is contained in its software rather than its hardware. The underlying chip itself, though by no means simple, is at its core a very large in-order processor. So it falls to the software stack to make Denver sing.

Accomplishing this task is NVIDIA's dynamic code optimizer (DCO). The purpose of the DCO is to accomplish two tasks: to translate ARM code to Denver's native format, and to optimize this code to make it run better on Denver. With no out-of-order hardware on Denver, it is the DCO's task to find instruction level parallelism within a thread to fill Denver's many execution units, and to reorder instructions around potential stalls, something that is no simple task.



The DCO system employed in the Denver CPU is codesigned software that extends ideas from prior system-level binary translators.² The primary function is to execute the user's code. The secondary function is to profile execution, create, optimize, and manage regions of tens to thousands of ARM instructions to form equivalent microcode-optimized regions that execute efficiently on the underlying microarchitecture.

More on NVIDIA Denver Code Optimizer

DENVER: NVIDIA'S FIRST 64-BIT ARM PROCESSOR

NVIDIA'S FIRST 64-BIT ARM PROCESSOR, CODE-NAMED DENVER, LEVERAGES A HOST OF NEW TECHNOLOGIES, SUCH AS DYNAMIC CODE OPTIMIZATION, TO ENABLE HIGH-PERFORMANCE MOBILE COMPUTING. IMPLEMENTED IN A 28-NM PROCESS, THE DENVER CPU CAN ATTAIN CLOCK SPEEDS OF UP TO 2.5 GHZ. THIS ARTICLE OUTLINES THE DENVER ARCHITECTURE, DESCRIBES ITS TECHNOLOGICAL INNOVATIONS, AND PROVIDES RELEVANT COMPARISONS AGAINST COMPETING MOBILE PROCESSORS.

Codesigning a hardware processor with a DCO software system creates both additional validation exposure and benefits. The DCO system can be upgraded in the field to address functionality, performance, or security issues.

The Denver hardware decoder provides a mechanism for periodically profiling recently taken branches. This branch history is moved into a shared buffer that can be processed from other cores, thereby minimizing the latency of the interruption. The DCO system will then run a thread that uses this profile to evaluate the dynamic properties of code executing and to assemble a picture of which code regions are hottest across all the processors. On finding sufficiently hot code, the DCO system will begin an optimization process to turn this input ARM code into a microcode execution region. The optimization process uses well-known traditional³ and more speculative compiler techniques to reduce work and increase efficiency of execution on the underlying skewed pipeline. To keep the latency of interruptions to a minimum, the optimizer thread is time-sliced with ARM execution (if any) and runs in a mode that can be quickly interrupted.

Transmeta: x86 to VLIW Translation




Figure 5. The Code Morphing software mediates between x86 software and the Crusoe processor.

Klaiber, "The Technology Behind Crusoe Processors," Transmeta White Paper 2000.

More on Static Instruction Scheduling

Trace Scheduling Idea

TRACE SCHEDULING LOOP-FREE CODE

43

58:18 / 1:41:17

Lecture 16. Static Instruction Scheduling - Carnegie Mellon - Comp. Arch. 2015 - Onur Mutlu

7,136 views • Feb 26, 2015

46 0 SHARE SAVE ...

Carnegie Mellon Computer Architecture
23K subscribers

Lecture 16: Static Instruction Scheduling
Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)
Date: Feb 23rd, 2015

Lecture 16 slides (pdf): <http://www.ece.cmu.edu/~ece447/s15/li...>

SUBSCRIBED

<https://www.youtube.com/watch?v=isBEVkljgGA&list=PL5PHm2jkkXmi5CxxI7b3JCL1TWybtDtKq&index=18>

Lectures on Static Instruction Scheduling

- **Computer Architecture, Spring 2015, Lecture 16**
 - Static Instruction Scheduling (CMU, Spring 2015)
 - <https://www.youtube.com/watch?v=isBEVkiJgGA&list=PL5PHm2jkkXmi5CxxI7b3JCL1TWybTDtKq&index=18>
- **Computer Architecture, Spring 2013, Lecture 21**
 - Static Instruction Scheduling (CMU, Spring 2013)
 - <https://www.youtube.com/watch?v=XdDUn2WtkRg&list=PL5PHm2jkkXmidJOd59REog9jDnPDTG6IJ&index=21>

Recall: How to Handle Data Dependences

- Anti and output dependences are easier to handle
 - ❑ write to the destination only in last stage and in program order
- Flow dependences are more interesting
- Six fundamental ways of handling flow dependences
 - ❑ Detect and wait until value is available in register file
 - ❑ Detect and forward/bypass data to dependent instruction
 - ❑ Detect and eliminate the dependence at the software level
 - No need for the hardware to detect dependence
 - ❑ Detect and move it out of the way for independent instructions
 - ❑ Predict the needed value(s), execute “speculatively”, and verify
 - ❑ Do something else (fine-grained multithreading)
 - No need to detect

Fine-Grained Multithreading

Fine-Grained Multithreading

- Idea: Fetch from a different thread every cycle such that no two instructions from a thread are in the pipeline concurrently
 - ❑ Hardware has multiple thread contexts (PC+registers per thread)
 - ❑ Threads are completely independent
 - ❑ No instruction is fetched from the same thread until the prior branch/instruction from the thread completes


+ No logic needed for handling control and data dependences within a thread

+ High thread-level throughput

-- Single thread performance suffers


-- Extra logic for keeping thread contexts

-- Throughput loss when there are not enough threads to keep the pipeline full



Each pipeline stage has an instruction from a different, completely-independent thread

Fine-Grained Multithreading: Basic Idea



Each pipeline stage has an instruction from a different, completely-independent thread

We need a PC and a register file for each thread + muxes and control

Fine-Grained Multithreading (II)

- Idea: Fetch from a different thread every cycle such that no two instructions from a thread are in the pipeline concurrently
- Tolerates control and data dependence resolution latencies by overlapping the latency with useful work from other threads
- Improves pipeline utilization by taking advantage of multiple threads
- Improves thread-level throughput but sacrifices per-thread throughput & latency
- Thornton, “Parallel Operation in the Control Data 6600,” AFIPS 1964.
- Smith, “A pipelined, shared resource MIMD computer,” ICPP 1978.

Fine-Grained Multithreading: History


- CDC 6600's peripheral processing unit is fine-grained multithreaded
 - ❑ Thornton, “[Parallel Operation in the Control Data 6600](#),” AFIPS 1964.
 - ❑ Processor executes a different I/O thread every cycle
 - ❑ An operation from the same thread is executed every 10 cycles
- Denelcor HEP (Heterogeneous Element Processor)
 - ❑ Smith, “[A pipelined, shared resource MIMD computer](#),” ICPP 1978.
 - ❑ 120 threads/processor
 - ❑ Available queue vs. unavailable (waiting) queue for threads
 - ❑ Each thread can have only 1 instruction in the processor pipeline
 - ❑ Each thread independent
 - ❑ To each thread, processor looks like a non-pipelined machine
 - ❑ [System throughput vs. single thread performance tradeoff](#)

Fine-Grained Multithreading in HEP


- Cycle time: 100ns
- 8 stages → 800 ns to complete an instruction
 - ❑ assuming no memory access
- No control and data dependence checking



Burton Smith
(1941-2018)



Multithreaded Pipeline Example



Sun Niagara Multithreaded Pipeline



Kongetira et al., "Niagara: A 32-Way Multithreaded Sparc Processor," IEEE Micro 2005.

Fine-Grained Multithreading

■ Advantages


- + No need for dependence checking between instructions
(only one instruction in pipeline from a single thread)
- + No need for branch prediction logic
- + Otherwise-bubble cycles used for executing useful instructions from different threads
- + Improved system throughput, latency tolerance, pipeline utilization

■ Disadvantages


- Extra hardware complexity: multiple hardware contexts (PCs, register files, ...), thread selection logic
- Reduced single thread performance (one instruction fetched every N cycles from the same thread)
- Resource contention between threads in caches and memory
- Dependence checking logic *between* threads may be needed (load/store)

Modern GPUs are FGMT Machines

NVIDIA GeForce GTX 285 “core”



64 KB of storage
for thread contexts
(registers)



= data-parallel (SIMD) func. unit,
control shared across 8 units

■ = multiply-add
■ = multiply




= instruction stream decode




= execution context storage

NVIDIA GeForce GTX 285 “core”



- Groups of 32 **threads** share instruction stream (each group is a Warp): they execute the same instruction on different data
- **Up to 32 warps are interleaved in an FGMT manner**
- Up to 1024 thread contexts can be stored

NVIDIA GeForce GTX 285 (~2009)



30 cores on the GTX 285: 30,720 threads

Further Reading for the Interested (I)

A PIPELINED, SHARED RESOURCE MIMD COMPUTER

Burton J. Smith
Denelcor, Inc.
Denver, Colorado 80205



Burton Smith
(1941-2018)

Architecture and applications of the HEP multiprocessor computer system

Burton J. Smith
Denelcor, Inc., 14221 E. 4th Avenue, Aurora, Colorado 80011

Further Reading for the Interested (II)

The Tera Computer System*

Robert Alverson

David Callahan
Allan Porterfield

Daniel Cummings
Burton Smith


Brian Koblenz

Tera Computer Company
Seattle, Washington USA

4 Processors

Each processor in a Tera computer can execute multiple instruction streams simultaneously. In the current implementation, as few as one or as many as 128 program counters may be active at once. On every tick of the clock, the processor logic selects a stream that is ready to execute and allows it to issue its next instruction. Since instruction interpretation is completely pipelined by the processor and by the network and memories as well, a new instruction from a different stream may be issued in each tick without interfering with its predecessors. When an instruction finishes, the stream to which it belongs thereby becomes ready to execute the next instruction. As long as there are enough instruction streams in the processor so that the average instruction latency is filled with instructions from other streams, the processor is being fully utilized. Thus, it is only necessary to have enough streams to hide the expected latency (perhaps 70 ticks on average); once latency is hidden the processor is running at peak performance and additional streams do not speed the result.

More on Multithreading (I)



The video player displays a lecture slide titled "Coarse-grained Multithreading". The slide content is as follows:

- Idea: When a thread is stalled due to some event, switch to a different hardware context
 - Switch-on-event multithreading

The video player interface includes a progress bar at 1:19:21 / 1:35:43, a like count of 10, and a comment count of 0. The video title is "Carnegie Mellon - Parallel Computer Architecture 2013 - Onur Mutlu - Lec 9 - Multithreading". The channel name is "Carnegie Mellon Computer Architecture" with 23K subscribers. The video description includes "Lecture 9: Multithreading", "Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)", and "Date: September 26, 2013.".

Carnegie Mellon - Parallel Computer Architecture 2013 - Onur Mutlu - Lec 9 - Multithreading

1,252 views • Nov 19, 2013

10 0 SHARE SAVE ...

 **Carnegie Mellon Computer Architecture**
23K subscribers

Lecture 9: Multithreading
Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)
Date: September 26, 2013.

ANALYTICS EDIT VIDEO

https://www.youtube.com/watch?v=iqi9wFqFiNU&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnyV6D&index=51

More on Multithreading (II)

Intel Pentium 4 Hyperthreading

Long latency load handling

- Multi-level scheduling window

More partitioned structures

- I-TLB
- Instruction Queues
- Store buffer
- Reorder buffer

5% area overhead due to SMT

Marr et al., "Hyper-Threading Technology Architecture and Microarchitecture," Intel Technology Journal 2002.

replay loop

4+st

split line fetch

Assigning Confidence to Conditional Branches

MICRO 1776

Carnegie Mellon -Parallel Computer Architecture 2012 - Onur Mutlu - Lecture 10 - Multithreading II

1,594 views • Sep 21, 2013

11 0 SHARE SAVE ...



Carnegie Mellon Computer Architecture
1.81K subscribers

Lecture 10: Multithreading II

Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)

Date: September 28, 2012.

SUBSCRIBED



https://www.youtube.com/watch?v=e8lfl6MbiLg&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnV6D&index=52

More on Multithreading (III)

g (Tandem, Compaq Himalaya)

Lockstepping (Tandem, Compaq Himalaya)

the processor, compare the results of two
re committing an instruction

the processor, compare the results of two
processors before committing an instruction

Check a function $f(x) = 1, x, x^2, x^3$

Carnegie Mellon - Parallel Computer Architecture 2013 - Onur Mutlu - Lec 13-Multi-threading II

1,132 views • Sep 21, 2013

8 0 SHARE SAVE ...



Carnegie Mellon Computer Architecture
1.81K subscribers

Lecture 13: Multi-threading III

Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)

Date: October 5, 2012.

SUBSCRIBED



https://www.youtube.com/watch?v=7vkDpZ1-hHM&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnyV6D&index=53

More on Multithreading (IV)



Carnegie Mellon - Parallel Computer Architecture 2013 - Onur Mutlu - Lec 15 - Speculation 1

915 views • Sep 21, 2013

9 0 SHARE SAVE ...



Carnegie Mellon Computer Architecture
1.81K subscribers

Lecture 15: Speculation I

Lecturer: Prof. Onur Mutlu (<http://users.ece.cmu.edu/~omutlu/>)

Date: October 10, 2012.

SUBSCRIBED



https://www.youtube.com/watch?v=-hbmzIDe0sA&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnnyV6D&index=54

Lectures on Multithreading

■ Parallel Computer Architecture, Fall 2012, Lecture 9

- ❑ Multithreading I (CMU, Fall 2012)
- ❑ https://www.youtube.com/watch?v=iqi9wFqFiNU&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnyV6D&index=51

■ Parallel Computer Architecture, Fall 2012, Lecture 10

- ❑ Multithreading II (CMU, Fall 2012)
- ❑ https://www.youtube.com/watch?v=e8lfl6MbILg&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnyV6D&index=52

■ Parallel Computer Architecture, Fall 2012, Lecture 13

- ❑ Multithreading III (CMU, Fall 2012)
- ❑ https://www.youtube.com/watch?v=7vkDpZ1-hHM&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnyV6D&index=53

■ Parallel Computer Architecture, Fall 2012, Lecture 15

- ❑ Speculation I (CMU, Fall 2012)
- ❑ https://www.youtube.com/watch?v=-hbmzIDe0sA&list=PL5PHm2jkkXmgDN1PLwOY_tGtUlynnyV6D&index=54