

Digital Design & Computer Arch.

Lecture 18b: Systolic Array Architectures

Prof. Onur Mutlu

ETH Zürich

Spring 2026

24 April 2026

Research Opportunities

- If you are interested in **doing research** in Comp. Architecture, Memory/Storage, Security, Systems & Bioinformatics:
 - **Email me** with your interest (CC: Konstantina & Ataberk)
 - Take **the seminar course and the “Computer Architecture” course**
 - **Do readings** and assignments on your own & **talk with us**
- There are **many exciting projects and research positions**, e.g.:
 - Novel memory/storage/computation/communication systems
 - New execution paradigms (e.g., in-memory computing)
 - Hardware security, safety, reliability, predictability
 - GPUs, TPUs, FPGAs, PIM, heterogeneous systems, ...
 - Security-architecture-reliability-energy-performance interactions
 - Architectures for genomics/proteomics/medical/health/AI/ML
 - A limited list is here: <https://safari.ethz.ch/theses/>

Bachelor's Seminar in Computer Architecture

- Fall 2026 (offered every Fall and Spring Semester)
- 2 credit units
- **Rigorous seminar on fundamental and cutting-edge topics in computer architecture**
- Critical paper presentation, review, and discussion of seminal and cutting-edge works in computer architecture
 - We will cover many ideas & issues, analyze their tradeoffs, perform **critical thinking** and **brainstorming**
- Participation, presentation, synthesis report, lots of discussion
- You can register for the course online
- https://safari.ethz.ch/architecture_seminar

Our Major Courses & Lectures

■ **First Digital Design & Computer Architecture Course (DDCA)**

- ❑ Freshman-level (CS & ECE): Digital Design and Computer Architecture
- ❑ Spring 2025 Livestream Edition:
<https://www.youtube.com/watch?v=ubhxKNIOIRg&list=PL5Q2soXY2Zi9Eo29LMgKVcaydS7V1zZW3&index=1>

■ **Fundamentals of Computer Architecture Course (FoCA)**

- ❑ Junior/Senior-level (ECE & CS): Computer Architecture
- ❑ Spring 2025 Livestream Edition:
https://www.youtube.com/watch?v=uKgMFj1eQQc&list=PL5Q2soXY2Zi_ZMtqz1r-GHm-zzuE1QfIg&index=1

■ **Advanced Computer Architecture Course (ACA)**

- ❑ Senior/Masters/PhD-level (CS & ECE): Computer Architecture
- ❑ Fall 2024 Livestream Edition:
<https://www.youtube.com/watch?v=ziMRjDLEwo&list=PL5Q2soXY2Zi-LfDdGgWyLcTSqzm6a26wD&index=1>

■ **Seminar in Computer Architecture (SCA)**

- ❑ Fall 2024 Livestream Edition:
<https://www.youtube.com/watch?v=SamRIUuT1FA&list=PL5Q2soXY2Zi-B6hFQA2y8lisISOYMH1&index=1>

Computer Architecture Course

- Computer Architecture Course (Fall 2024):
 - ❑ <https://safari.ethz.ch/architecture/fall2024>
 - ❑ <https://www.youtube.com/onurmutlulectures>
 - ❑ <https://www.youtube.com/watch?v=ziMRjDILEwo&list=PL5Q2soXY2Zi-LfDdGgWyLcTSqzm6a26wD&index=1>

The screenshot shows a Zoom video player interface. The main video area displays a slide with the title "Computer Architecture Lecture 1: Introduction and Basics" in red and green text. Below the title, it says "Prof. Onur Mutlu", "ETH Zürich", "Fall 2024", and "19 September 2024". The Zoom logo is visible in the bottom right corner of the video area. On the right side of the video player, there is a list of 13 lectures, each with a thumbnail and a duration. The lectures are:

- 1. Computer Architecture - Lecture 1: Introduction and... (2:43:48)
- 2. Computer Architecture - Lecture 2: Memory Systems... (2:34:02)
- 3. Computer Architecture - Lecture 3: Memory-Centric... (2:43:20)
- 4. Computer Architecture - Lecture 4: Memory Centric... (2:30:54)
- 5. Computer Architecture - Lecture 5: Memory... (2:48:05)
- 6. Computer Architecture - Lecture 6: Cutting-Edge... (2:39:41)
- 7. Computer Architecture - Lecture 7: RowHammer... (2:42:48)
- 8. Computer Architecture - Lecture 8: Data Retention in... (2:46:44)
- 9. Computer Architecture - Lecture 9: Memory Latency II... (2:44:40)
- 10. Computer Architecture - Lecture 10: Memory... (2:30:54)
- 11. Computer Architecture - Lecture 11: Memory... (2:47:01)
- 12. Computer Architecture - Lecture 12: Memory-Centric... (3:04:42)
- 13. Computer Architecture - Lecture 13: Memory... (2:36:01)

Computer Architecture - Lecture 1: Introduction and Basics (Fall 2024)

Onur Mutlu Lectures
55.3K subscribers

Analytics

Edit video

28,859 views Streamed live on Sep 19, 2024
Computer Architecture, ETH Zürich, Fall 2024 (<https://safari.ethz.ch/architecture/f...>)

544 Like Share Promote Download

Top chat replay

Top Fans

For More Information: SAFARI Website



Think BIG, Aim HIGH!

<https://safari.ethz.ch>

<https://people.inf.ethz.ch/omutlu/projects.htm>

Suggestions on Education, Research, Growth



The video player shows a presentation slide with the following content:

Applying to Grad School
& Doing Impactful Research

Onur Mutlu
omutlu@gmail.com
<https://people.inf.ethz.ch/omutlu>
13 June 2020
Undergraduate Architecture Mentoring Workshop @ ISCA 2021


Logos for SAFARI, ETH zürich, and Carnegie Mellon are displayed at the bottom of the slide.

Below the video player, the YouTube interface shows the video title: "Arch. Mentoring Workshop @ISCA'21 - Applying to Grad School & Doing Impactful Research - Onur Mutlu". It also displays "1,563 views • Premiered Jun 16, 2021". The channel name "Onur Mutlu Lectures" with "17.2K subscribers" is visible. Interaction buttons for likes (74), comments (1), share, save, and analytics/edit video are present. A description of the panel talk at the Undergraduate Architecture Mentoring Workshop at ISCA 2021 is provided at the bottom.

Onur Mutlu's SAFARI Research Group

Computer architecture, HW/SW, systems, bioinformatics, security, memory

<https://safari.ethz.ch/safari-newsletter-january-2021/>



SAFARI
SAFARI Research Group
safari.ethz.ch

Think BIG, Aim HIGH!

SAFARI

<https://safari.ethz.ch>

SAFARI Newsletter January 2021 Edition

- <https://safari.ethz.ch/safari-newsletter-january-2021/>



SAFARI
SAFARI Research Group

Newsletter
January 2021

*Think Big, Aim High, and
Have a Wonderful 2021!*



Dear SAFARI friends,

Happy New Year! We are excited to share our group highlights with you in this second edition of the SAFARI newsletter (You can find the first edition from April 2020 [here](#)). 2020 has

SAFARI Newsletter December 2021 Edition

- <https://safari.ethz.ch/safari-newsletter-december-2021/>

SAFARI
SAFARI Research Group

Think Big, Aim High

ETH zürich



View in your browser
December 2021



SAFARI Newsletter June 2023 Edition

- <https://safari.ethz.ch/safari-newsletter-june-2023/>

SAFARI
SAFARI Research Group

Think Big, Aim High

ETH zürich



View in your browser

June 2023



SAFARI Newsletter July 2024 Edition

- <https://safari.ethz.ch/safari-newsletter-july-2024/>



SAFARI Newsletter December 2025 Edition

- <https://people.inf.ethz.ch/omutlu/pub/safari-newsletter-december-2025.pdf>

1st SAFARI Conference 2025



Abdullah Giray
Yağlıkçı, CISPA



Can Firtina
University of
Maryland, College
Park



Haiyu Mao
King's College
London



Juan Gómez
Luna, NVIDIA



Christina Giannoula
MPI-SWS



Jawad Haj-
Yahya, Rivos Inc.



Nastaran
Hajinazar
Intel Labs



Lois Orosa
Galicia Supercomputing
Center, CESGA



Yu (Lenny) Liang
Inria Paris



Gagandeep
Singh, AMD



Minesh Patel
Rutgers University



Geraldo Francisco
de Oliveira
ETH Zurich



Saugata Ghose
University of Illinois
Urbana-Champaign

Monday, Dec 15 2025, 8:30-18:30 CET

ETH zürich

SAFARI

SAFARI Introduction & Research

Computer architecture, HW/SW, systems, bioinformatics, security, memory



Seminar in Computer Architecture - Lecture 5: Potpourri of Research Topics (Spring 2023)



Onur Mutlu Lectures
32.6K subscribers



719 views Streamed 1 month ago Livestream - Seminar in Computer Architecture - ETH Zürich (Spring 2023)



Think BIG, Aim HIGH!

SAFARI

<https://www.youtube.com/watch?v=mV2OuB2djEs>

SAFARI PhD and Post-Doc Alumni

- Nika Mansouri Ghiasi (ETH Zurich)
- Rahul Bera (ETH Zurich), **MICRO 2022 Best Paper Award, ISCA 2024 Best Paper Award**
- Geraldo Francisco de Oliveira Jr. (ETH Zurich)
- Can Firtina (UMD, Assistant Professor), **ETH Doctoral Medal 2025**
- Lukas Breitwieser (CERN), **2025 Impact Award of the Research Software Engineering Society, Best artifact award at PPOPP 2023**
- A. Giray Yaglikci (CISPA), **2025 SIGMICRO Dissertation Award, DSN Carter Award Best Thesis 2025; IEEE TCCA/ACM SIGARCH Dissertation Award HM 2025, PACT 2023 SRC Winner, Intel HWSec Award Finalist 2021, HOST PhD Competition Finalist 2024, VTS PhD Competition Finalist 2025, HPCA HoF 2025, ISCA HoF 2025, Best Artifact award at HPCA 2025/ISCA 2023...**
- Hasan Hassan (Rivos), **EDAA Outstanding Dissertation Award 2023; S&P 2020 Best Paper Award, 2020 Pwnie Award, IEEE Micro TP HM 2020**
- Christina Giannoula (MPI-SWS), **NTUA Best Dissertation Award 2023**
- Minesh Patel (Rutgers, Asst. Prof.), **DSN Carter Award Best Thesis 2022; ETH Doctoral Medal 2023; MICRO'20 & DSN'19 Best Paper Awards; ISCA HoF 2021**
- Damla Senol Cali (Bionano Genomics), **SRC TECHCON 2019 Best Student Presentation Award; RECOMB-Seq 2018 Best Poster Award**
- Nastaran Hajinazar (Intel)
- Gagandeep Singh (AMD/Xilinx), **FPL 2020 Best Paper Award Finalist**
- Amirali Boroumand (Stanford Univ → Google), **SRC TECHCON 2018 Best Presentation Award**
- Jeremie Kim (Apple), **EDAA Outstanding Dissertation Award 2020; IEEE Micro Top Picks 2019; ISCA/MICRO HoF 2021**
- Nandita Vijaykumar (Univ. of Toronto, Assistant Professor), **ISCA Hall of Fame 2021**
- Kevin Hsieh (Microsoft Research, Principal Researcher)
- Justin Meza (Facebook), **HiPEAC 2015 Best Student Presentation Award; ICCD 2012 Best Paper Award**
- Mohammed Alser (Georgia State Univ., Assistant Professor), **IEEE Turkey Best PhD Thesis Award 2018**
- Yixin Luo (Google → OpenAI), **HPCA 2015 Best Paper Session**
- Kevin Chang (Facebook), **SRC TECHCON 2016 Best Student Presentation Award**
- Rachata Ausavarungnirun (KMUTB, Assistant Professor → MangoBoost), **NOCS 2015 and NOCS 2012 Best Paper Award Finalist**
- Gennady Pekhimenko (Univ. of Toronto, Associate Professor), **ISCA Hall of Fame 2021; ASPLOS 2015 SRC Winner**
- Vivek Seshadri (Microsoft Research, Principal Researcher)
- Donghyuk Lee (NVIDIA Research, Senior Researcher), **HPCA Hall of Fame 2018**
- Yoongu Kim (Software Robotics → Google), **IFIP JCL Award'24, TCAD'19 Top Pick Award; IEEE Micro Top Picks'10; HPCA'10 Best Paper Session**
- Lavanya Subramanian (Intel Labs → Facebook → OpenAI)
- Samira Khan (Univ. of Virginia, Assistant Professor → Google), **HPCA 2014 Best Paper Session, DSN 2025 Test of Time Award**
- Saugata Ghose (Univ. of Illinois, Assistant Professor), **DFRWS-EU 2017 Best Paper Award, HPCA HoF 2024, MICRO HoF 2024**
- Jawad Haj-Yahya (Huawei Research Zurich, Principal Researcher → Rivos)
- Lois Orosa (Galicia Supercomputing Center, Director)
- Jisung Park (POSTECH, Assistant Professor)
- Gagandeep Singh (AMD/Xilinx, Researcher)
- Juan Gomez-Luna (NVIDIA, Researcher), **ISPASS 2023 Best Paper Session, MICRO HoF 2022**
- Mohammed Alser (Georgia State Univ., Assistant Professor), **IEEE Turkey Best PhD Thesis Award 2018**
- Haiyu Mao (King's College, Assistant Professor)
- Yu (Lenny) Liang (INRIA Paris, Researcher)

You Can Join Us!

■ <https://safari.ethz.ch/apply/>

SAFARI Researcher Applications

Sign in

This is the application submission site to be considered for being a researcher in the [SAFARI Research Group](#), directed by [Professor Onur Mutlu \(Publications and Teaching\)](#).

If you are interested in doing research in the [SAFARI Research Group](#), please make sure you apply through this submissions site and supply as many of the requested documents and information as possible. Please read and follow the provided instructions and submit as complete an application as possible (given the position you are applying for).

We suggest studying the following materials before submission:

[SAFARI Publications and Courses](#)

[Onur Mutlu's Online Lectures and Course Materials](#)

We strongly recommend that you read and analyze critically as many recent papers from our group as possible. This is the best way to prepare for the application process. Our recommendation is that you use professor Mutlu's methodology for critically analyzing papers.

[Guide On Reviewing Papers](#)

Good luck!

Welcome to the SAFARI at ETH Zurich -- PhD, Postdoc, Internship, Visiting Researcher Applications (SAFARI Researcher Applications) submissions site.

Approaches to (Instruction-Level) Concurrency

- Pipelining
- Fine-Grained Multithreading
- Out-of-order Execution
- Dataflow (at the ISA level)
- Superscalar Execution
- VLIW
- Systolic Arrays
- Decoupled Access Execute
- SIMD Processing (Vector and array processors, GPUs)

Readings for Today

■ Required

- H. T. Kung, “Why Systolic Architectures?,” IEEE Computer 1982.

■ Recommended

- Jouppi et al., “In-Datacenter Performance Analysis of a Tensor Processing Unit,” ISCA 2017.
- Jouppi et al., “A Domain-Specific Supercomputer for Training Deep Neural Networks,” CACM 2020.

Readings for Next Week

■ Required

- Lindholm et al., "NVIDIA Tesla: A Unified Graphics and Computing Architecture," IEEE Micro 2008.

■ Recommended


- Peleg and Weiser, "MMX Technology Extension to the Intel Architecture," IEEE Micro 1996.

Systolic Arrays

General Purpose vs. Special Purpose Systems


General Purpose

CPU_s




Apple M1

GPU_s



Nvidia GTX 1070


FPGAs



Xilinx Spartan

Special Purpose

ASIC_s



Cerebras WSE-2



Flexible: Can execute any program

Easy to program & use

Not the best performance & efficiency

Efficient & High performance

(Usually) Difficult to program & use

Inflexible: Limited set of programs

Systolic Arrays: Specialized Accelerators

- **Systolic arrays comprise an execution model**
 - They implement systolic computation
 - Different from **von Neumann**, different from **dataflow**
- They are initially designed as **special-purpose accelerators**
 - For convolutions, filtering, pattern matching, special-purpose matrix/vector computations in image & vision processing, signal processing, pattern recognition, etc.
- They are currently used heavily in machine learning
 - Specialized machine learning accelerators
- Their general execution model can be (and is) generalized
 - As we will see later in this lecture

Systolic Arrays: Motivation

- Goal: design an accelerator that has
 - Simple, regular design (keep # unique parts small and regular)
 - High concurrency → high performance
 - Balanced computation and I/O (memory) bandwidth
- Idea: Replace a single processing element (PE) with a regular array of PEs and carefully orchestrate flow of data between the PEs
 - such that they collectively transform a piece of input data before outputting it to memory
- Benefit: Maximizes computation done on a single piece of data element brought from memory

Systolic Arrays




Figure 1. Basic principle of a systolic system.

Memory: heart
Data: blood
PEs: cells

Memory pulses
data through
PEs

- H. T. Kung, "[Why Systolic Architectures?](#)," IEEE Computer 1982.

Why Systolic Architectures?

- Idea: Data flows from the computer memory in a rhythmic fashion, passing through many processing elements before it returns to memory
- Similar to blood flow: heart → many cells → heart
 - Different cells “process” the blood
 - Many veins operate simultaneously
 - Can be many-dimensional
- Why? Special purpose accelerators/architectures need
 - Simple, regular design (keep # unique parts small and regular)
 - High concurrency → high performance
 - Balanced computation and I/O (memory) bandwidth

Systolic Architectures

- Basic principle: Replace a single PE with a **regular array of PEs** and **carefully orchestrate flow of data** between the PEs
 - Balance computation and memory bandwidth




Figure 1. Basic principle of a systolic system.

- Differences from pipelining:
 - These are individual PEs
 - Array structure can be non-linear and multi-dimensional
 - PE connections can be multidirectional (and different speed)
 - PEs can have local memory and execute kernels (rather than a piece of the instruction)

Systolic Computation Example

■ Convolution

- Used in filtering, pattern matching, correlation, polynomial evaluation, etc ...
- Many **image processing** tasks
- **Machine learning**: up to hundreds of **convolutional layers** in Convolutional Neural Networks (CNN)


Given the sequence of weights $\{w_1, w_2, \dots, w_k\}$
and the input sequence $\{x_1, x_2, \dots, x_n\}$,

compute the result sequence $\{y_1, y_2, \dots, y_{n+1-k}\}$
defined by

$$y_i = w_1x_i + w_2x_{i+1} + \dots + w_kx_{i+k-1}$$


LeNet-5, a Convolutional Neural Network for Hand-Written Digit Recognition

This is a 1024*8 bit input, which will have a truth table of 2^{8196} entries



An Example of 2D Convolution

Output feature map



Input feature map

Structure information

Input: 5*5 (blue)

Kernel (filter): 3*3 (grey)

Output: 5*5 (green)

Computation information

Stride: 1

Padding: 1 (white)

Output Dim = (Input + 2*Padding
- Kernel) / Stride + 1

An Example of 2D Convolution

**Input
Layer**

0	2	5	-3	10
9	-1	29	4	11
0	34	7	-9	-17
6	25	6	0	0
-8	21	0	77	0

**CNN
kernel**

3	2	1
-1	-2	-3
2	1	-1

**Output
Layer**

-58	

Convolutional Neural Networks: Demo

[Back to Yann's Home Publications](#)

LeNet-5 Demos

Unusual Patterns

[unusual styles](#)
[weirdos](#)

Invariance

[translation](#) (anim)
[scale](#) (anim)
[rotation](#) (anim)
[squeezing](#) (anim)
[stroke width](#) (anim)

Noise Resistance

[noisy 3 and 6](#)
[noisy 2](#) (anim)
[noisy 4](#) (anim)

Multiple Character

[various stills](#)
[dancing 00](#) (anim)
[dancing 384](#) (anim)

Complex cases

(anim)
[35 -> 53](#)
[12 -> 4 -> 21](#)
[23 -> 32](#)
[30 + noise](#)
[31-51-57-61](#)


LeNet-5, convolutional neural networks

Convolutional Neural Networks are a special kind of multi-layer neural networks. Like almost every other neural networks they are trained with a version of the back-propagation algorithm. Where they differ is in the architecture.

Convolutional Neural Networks are designed to recognize visual patterns directly from pixel images with minimal preprocessing.

They can recognize patterns with extreme variability (such as handwritten characters), and with robustness to distortions and simple geometric transformations.

LeNet-5 is our latest convolutional network designed for handwritten and machine-printed character recognition. Here is an example of LeNet-5 in action.



Many more examples are available in the column on the left:


Several papers on LeNet and convolutional networks are available on my [publication page](#):

[LeCun et al., 1998]

Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner.
Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, november 1998.
[ps.gz](#)


[Bottou et al., 1997]

L. Bottou, Y. LeCun, and Y. Bengio. Global training of



<http://yann.lecun.com/exdb/lenet/index.html>

Implementing a Convolutional Layer with Matrix Multiplication



$$\begin{array}{|c|c|c|c|c|c|c|c|c|c|c|c|} \hline 1 & 1 & 2 & 2 & 1 & 1 & 1 & 1 & 0 & 1 & 1 & 0 \\ \hline 1 & 0 & 0 & 1 & 2 & 1 & 2 & 1 & 1 & 2 & 2 & 0 \\ \hline \end{array} * \begin{array}{|c|c|c|c|} \hline 1 & 2 & 1 & 1 \\ \hline 2 & 0 & 1 & 3 \\ \hline 1 & 1 & 0 & 2 \\ \hline 1 & 3 & 2 & 2 \\ \hline 0 & 2 & 0 & 1 \\ \hline 2 & 1 & 1 & 2 \\ \hline 0 & 1 & 1 & 1 \\ \hline 1 & 2 & 1 & 0 \\ \hline 1 & 2 & 0 & 1 \\ \hline 2 & 1 & 1 & 3 \\ \hline 0 & 1 & 3 & 3 \\ \hline 1 & 3 & 3 & 2 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 12 & 18 & 13 & 22 \\ \hline 10 & 20 & 15 & 22 \\ \hline \end{array}$$

Convolution Filters W'

Input Features X (unrolled)

Output Features Y

Power of **Convolutions** and **Applied Courses**

- In 2010, Prof. Andreas Moshovos adopted Professor Hwu's ECE498AL Programming Massively Parallel Processors Class
- Several of Prof. Geoffrey Hinton's graduate students took the course
- These students developed the GPU implementation of the Deep CNN that was trained with 1.2M images to win the ImageNet competition

Example: AlexNet (2012)

- AlexNet wins the **ImageNet classification competition** with ~10% points higher accuracy than state-of-the-art
 - Krizhevsky et al., “**ImageNet Classification with Deep Convolutional Neural Networks**”, NIPS 2012.

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

Abstract

We trained a large, deep convolutional neural network to classify the 1.2 million high-resolution images in the ImageNet LSVRC-2010 contest into the 1000 different classes. On the test data, we achieved top-1 and top-5 error rates of 37.5% and 17.0% which is considerably better than the previous state-of-the-art. The neural network, which has 60 million parameters and 650,000 neurons, consists of five convolutional layers, some of which are followed by max-pooling layers, and three fully-connected layers with a final 1000-way softmax. To make training faster, we used non-saturating neurons and a very efficient GPU implementation of the convolution operation. To reduce overfitting in the fully-connected layers we employed a recently-developed regularization method called “dropout” that proved to be very effective. We also entered a variant of this model in the ILSVRC-2012 competition and achieved a winning top-5 test error rate of 15.3%, compared to 26.2% achieved by the second-best entry.

Example: GoogLeNet (2014)

- Google improves accuracy by **adding more network layers**
 - From 8 in AlexNet to 22 in GoogLeNet
 - Szegedy et al., “**Going Deeper with Convolutions**”, CVPR 2015.

Going Deeper with Convolutions

Christian Szegedy¹, Wei Liu², Yangqing Jia¹, Pierre Sermanet¹, Scott Reed³,
Dragomir Anguelov¹, Dumitru Erhan¹, Vincent Vanhoucke¹, Andrew Rabinovich⁴

¹Google Inc. ²University of North Carolina, Chapel Hill

³University of Michigan, Ann Arbor ⁴Magic Leap Inc.

¹{szegedy, jia, yq, sermanet, dragomir, dimitru, vanhoucke}@google.com

²wliu@cs.unc.edu, ³reedscott@umich.edu, ⁴arabinovich@magic Leap Inc.


Example: ResNet (2015)

- He et al., “Deep Residual Learning for Image Recognition”, CVPR 2016.

Deep Residual Learning for Image Recognition


Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun
Microsoft Research
{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

ImageNet experiments




Neural Network Layer Examples

LeNet



AlexNet



Systolic Computation Example: Convolution (I)

■ Convolution

- ❑ Used in filtering, pattern matching, correlation, polynomial evaluation, etc ...
- ❑ Many **image processing** tasks
- ❑ **Machine learning**: up to hundreds of **convolutional layers** in Convolutional Neural Networks (CNN)

Given the sequence of weights $\{w_1, w_2, \dots, w_k\}$
and the input sequence $\{x_1, x_2, \dots, x_n\}$,

compute the result sequence $\{y_1, y_2, \dots, y_{n+1-k}\}$
defined by

$$y_i = w_1x_i + w_2x_{i+1} + \dots + w_kx_{i+k-1}$$

Systolic Computation Example: Convolution (II)

- $y_1 = w_1x_1 + w_2x_2 + w_3x_3$
- $y_2 = w_1x_2 + w_2x_3 + w_3x_4$
- $y_3 = w_1x_3 + w_2x_4 + w_3x_5$




Figure 8. Design W1: systolic convolution array (a) and cell (b) where w_i 's stay and x_i 's and y_i 's move systolically in opposite directions.

Systolic Computation Example: Convolution (III)




Figure 10. Overlapping the executions of multiply and add in design W1.

- Worthwhile to implement adder and multiplier separately to allow overlapping of add/mul executions

Systolic Computation Example: Convolution (IV)

- One needs to **carefully orchestrate** when **data elements are input to the array**
- And when **output is buffered**
- This gets more involved when
 - Array dimensionality increases
 - PEs are less predictable in terms of latency

Example 2D Systolic Array Computation

- Multiply two 3x3 matrices (inputs)
 - Keep the final result in PE accumulators

$$\begin{bmatrix} c_{00} & c_{01} & c_{02} \\ c_{10} & c_{11} & c_{12} \\ c_{20} & c_{21} & c_{22} \end{bmatrix} = \begin{bmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{bmatrix} \times \begin{bmatrix} b_{00} & b_{01} & b_{02} \\ b_{10} & b_{11} & b_{12} \\ b_{20} & b_{21} & b_{22} \end{bmatrix}$$





Figure 1: A systolic array processing element

$$\begin{aligned} P &= M \\ Q &= N \\ R &= R + M * N \end{aligned}$$



Two-Dimensional Systolic Arrays




Figure 11. Two-dimensional systolic arrays: (a) type R, (b) type H, and (c) type T.

To a given problem there could be both one- and two-dimensional systolic array solutions. For example, two-dimensional convolution can be performed by a one-dimensional systolic array^{24,25} or a two-dimensional systolic array.⁶ When the memory speed is more than cell speed, two-dimensional systolic arrays such as those depicted in Figure 11 should be used. At each cell cycle, all the I/O ports on the array boundaries can input or output data items to or from the memory; as a result, the available memory bandwidth can be fully utilized. Thus, the choice of a one- or two-dimensional scheme is very dependent on how cells and memories will be implemented.

Combinations

- Systolic arrays can be chained together to form powerful systems
- This systolic array is capable of producing on-the-fly least-squares fit to all the data that has arrived up to any given moment




Figure 12. On-the-fly least-squares solutions using one- and two-dimensional systolic arrays, with $p = 4$.

Systolic Arrays: Pros and Cons

■ Advantages:

- ❑ **Principled**: Efficiently makes use of limited memory bandwidth, balances computation to I/O bandwidth availability
- ❑ **Specialized** (computation should fit the PE organization/functions)
 - improved **efficiency, simple design, high concurrency/performance**
 - good to do **more with less memory bandwidth** requirement

■ Downside:

- ❑ **Specialized**
 - **not generally applicable** because computation needs to fit the PE functions/organization

More Programmability in Systolic Arrays

- Each PE in a systolic array
 - Can store multiple “weights”
 - Weights can be selected on the fly
 - Eases implementation of, e.g., adaptive filtering
- Taken further
 - Each PE can have its own data and instruction memory
 - Data memory → to store partial/temporary results, constants
 - Leads to **stream processing, pipeline parallelism**
 - More generally, **staged execution**

Pipeline-Parallel (Pipelined) Programs





Figure 1. (a) The code of a loop, (b) Each iteration is split into 3 pipeline stages: A, B, and C. Iteration i comprises A_i , B_i , C_i . (c) Sequential execution of 4 iterations. (d) Parallel execution of 6 iterations using pipeline parallelism on a three-core machine. Each stage executes on one core.

Stages of Pipelined Programs

- Loop iterations are divided into **code segments called stages**
- Threads execute stages on different cores
 - Cores can be specialized for the computations that execute on them

```
loop {  
  Compute1  A  
  Compute2  B  
  Compute3  C  
}
```



Pipelined File Compression Example




Figure 3. File compression algorithm executed using pipeline parallelism

Systolic Array: Advantages & Disadvantages

■ Advantages

- ❑ Special purpose → high efficiency
- ❑ Makes multiple uses of each data item → reduced need for fetching/refetching → better use of memory bandwidth
- ❑ High concurrency
- ❑ Regular design (both data and control flow) → easier to implement

■ Disadvantages

- ❑ Not good at exploiting irregular parallelism
- ❑ Special purpose → not generally applicable
 - Needs software, programmer support to become more general purpose
- ❑ Difficult to program if problem does not fit (well)

Example Systolic Array: The WARP Computer

- HT Kung, CMU, 1984-1988
- Linear array of 10 cells, each cell a 10 Mflop programmable processor
- Attached to a general purpose host machine
- HLL and optimizing compiler to program the systolic array
- Used extensively to accelerate vision and robotics tasks
- Annaratone et al., “Warp Architecture and Implementation,” ISCA 1986.
- Annaratone et al., “The Warp Computer: Architecture, Implementation, and Performance,” IEEE TC 1987.

The WARP Computer




Figure 1: Warp system overview

The WARP Cell




Figure 2: Warp cell data path

An Example Modern Systolic Array: TPU (I)




Figure 3. TPU Printed Circuit Board. It can be inserted in the slot for an SATA disk in a server, but the card uses PCIe Gen3 x16.





Figure 4. Systolic data flow of the Matrix Multiply Unit. Software has the illusion that each 256B input is read at once, and they instantly update one location of each of 256 accumulator RAMs.

Jouppi et al., “In-Datacenter Performance Analysis of a Tensor Processing Unit”, ISCA 2017.

An Example Modern Systolic Array: TPU (II)

As reading a large SRAM uses much more power than arithmetic, the matrix unit uses systolic execution to save energy by reducing reads and writes of the Unified Buffer [Kun80][Ram91][Ovt15b]. Figure 4 shows that data flows in from the left, and the weights are loaded from the top. A given 256-element multiply-accumulate operation moves through the matrix as a diagonal wavefront. The weights are preloaded, and take effect with the advancing wave alongside the first data of a new block. Control and data are pipelined to give the illusion that the 256 inputs are read at once, and that they instantly update one location of each of 256 accumulators. From a correctness perspective, software is unaware of the systolic nature of the matrix unit, but for performance, it does worry about the latency of the unit.



Jouppi et al., “In-Datacenter Performance Analysis of a Tensor Processing Unit”, ISCA 2017.

Recall: Example 2D Systolic Array Computation

- Multiply two 3x3 matrices (inputs)
 - Keep the final result in PE accumulators

$$\begin{bmatrix} c_{00} & c_{01} & c_{02} \\ c_{10} & c_{11} & c_{12} \\ c_{20} & c_{21} & c_{22} \end{bmatrix} = \begin{bmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{bmatrix} \times \begin{bmatrix} b_{00} & b_{01} & b_{02} \\ b_{10} & b_{11} & b_{12} \\ b_{20} & b_{21} & b_{22} \end{bmatrix}$$





Figure 1: A systolic array processing element

$$\begin{aligned} P &= M \\ Q &= N \\ R &= R + M * N \end{aligned}$$



An Example Modern Systolic Array: TPU (III)




Figure 1. TPU Block Diagram. The main computation part is the yellow Matrix Multiply unit in the upper right hand corner. Its inputs are the blue Weight FIFO and the blue Unified Buffer (UB) and its output is the blue Accumulators (Acc). The yellow Activation Unit performs the nonlinear functions on the Acc, which go to the UB.

An Example Modern Systolic Array: TPU2



<https://www.nextplatform.com/2017/05/17/first-depth-look-googles-new-second-generation-tpu/>

4 TPU chips
vs 1 chip in TPU1

High Bandwidth Memory
vs DDR3

Floating point operations
vs FP16

45 TFLOPS per chip
vs 23 TOPS

Designed for training
and inference
vs only inference

TPU v2

Figure 1. A 2D-torus topology. TPUv2 uses a 16x16 2D torus.





Figure 2. Block diagram of a TensorCore (our internal development name for a TPU core, and not related to the Tensor Cores of NVIDIA GPUs).




3. The *Core Sequencer* fetches *VLIW* (*Very Long Instruction Word*) instructions from the core's on-chip, software-managed *Instruction Memory* (*Imem*), executes scalar operations using a 4K 32-bit scalar data memory (*Smem*) and 32 32-bit scalar registers (*Sregs*), and forwards vector instructions to the VPU. The 322-bit VLIW instruction can launch eight operations: two scalar, two vector ALU, vector load and store, and a pair of slots that queue data to and from the matrix

multiply and transpose units. The XLA compiler schedules loading *Imem* via independent overlays of code, as unlike conventional CPUs, there is no instruction cache.

4. The *Vector Processing Unit* (*VPU*) performs vector operations using a large on-chip *vector memory* (*Vmem*) with 32K 128 x 32-bit elements (16MiB), and 32 2D *vector registers* (*Vregs*) that each contain 128 x 8 32-bit elements (4 KiB). The VPU streams data to and from the MXU through decoupling FIFOs. The VPU collects and distributes data to *Vmem* via *data-level parallelism* (2D matrix and vector functional units) and *instruction-level parallelism* (8 operations per instruction).

An Example Modern Systolic Array: TPU3




32GB HBM per chip
vs 16GB HBM in TPU2

4 Matrix Units per chip
vs 2 Matrix Units in TPU2

90 TFLOPS per chip
vs 45 TFLOPS in TPU2


An Example Modern Systolic Array: TPU4



New ML applications (vs. TPU3):

- Computer vision
- Natural Language Processing (NLP)
- Recommender system
- Reinforcement learning that plays Go

250 TFLOPS per chip in 2021
vs 90 TFLOPS in TPU3



1 ExaFLOPS per board

<https://spectrum.ieee.org/tech-talk/computing/hardware/heres-how-googles-tpu-v4-ai-chip-stacked-up-in-training-tests>

A Summary Reading of Google TPUs

Ten Lessons From Three Generations Shaped Google's TPUv4i Industrial Product

Norman P. Jouppi, Doe Hyun Yoon, Matthew Ashcraft, Mark Gottscho, Thomas B. Jablin, George Kurian, James Laudon, Sheng Li, Peter Ma, Xiaoyu Ma, Thomas Norrie, Nishant Patil, Sushma Prasad, Cliff Young, Zongwei Zhou, and David Patterson, Google LLC

A Recent Reading on the TPUv4

TPU v4: An Optically Reconfigurable Supercomputer for Machine Learning with Hardware Support for Embeddings **Industrial Product***

Norman P. Jouppi, George Kurian, Sheng Li, Peter Ma, Rahul Nagarajan, Lifeng Nai, Nishant Patil, Suvinay Subramanian, Andy Swing, Brian Towles, Cliff Young, Xiang Zhou, Zongwei Zhou, and David Patterson
Google, Mountain View, CA

TPU Generations

TPU AI accelerators

TPU v1

2015
Internal inference
accelerator



TPU v2

2018
Domain-specific AI
supercomputing
256 chips
distributed shared
memory




TPU v3

2020
Liquid cooled
1k chips distributed
shared memory



TPU v4

2022
Optically
reconfigurable
3D Torus
4k chips with
distributed shared
memory



TPU v5e

2023
Purpose-built
cost-efficiency and
performance for
medium/large-scale
training and
inference



TPU v5p

2023
Our powerful,
scalable, and
flexible AI
accelerator



Trillium

2024
Designed for
exceptional
performance and
efficiency. Enabling
the next frontier of
AI models



An Analysis of the Google Edge TPU

- Amirali Boroumand, Saugata Ghose, Berkin Akin, Ravi Narayanaswami, Geraldo F. Oliveira, Xiaoyu Ma, Eric Shiu, and Onur Mutlu,
["Google Neural Network Models for Edge Devices: Analyzing and Mitigating Machine Learning Inference Bottlenecks"](#)
Proceedings of the 30th International Conference on Parallel Architectures and Compilation Techniques (PACT), Virtual, September 2021.
[[Slides \(pptx\)](#)] [[pdf](#)]
[[Talk Video](#) (14 minutes)]

**> 90% of the total system energy
is spent on **memory** in large ML models**

Google Neural Network Models for Edge Devices: Analyzing and Mitigating Machine Learning Inference Bottlenecks

Amirali Boroumand^{†◇}
Geraldo F. Oliveira^{*}


Saugata Ghose[‡]
Xiaoyu Ma[§]

Berkin Akin[§]
Eric Shiu[§]

Ravi Narayanaswami[§]
Onur Mutlu^{*†}

[†]Carnegie Mellon Univ. [◇]Stanford Univ. [‡]Univ. of Illinois Urbana-Champaign [§]Google ^{*}ETH Zürich


Energy Wasted on Data Movement



**In LSTMs and Transducers used by Google,
>90% energy spent on off-chip interconnect and DRAM**

Other Machine Learning Accelerators

Cerebras's Wafer Scale Engine (2019)



Cerebras WSE

1.2 Trillion transistors

46,225 mm²

- The largest ML accelerator chip
- 400,000 cores



Largest GPU

21.1 Billion transistors


815 mm²

NVIDIA TITAN V

<https://www.anandtech.com/show/14758/hot-chips-31-live-blogs-cerebras-wafer-scale-deep-learning>

<https://www.cerebras.net/cerebras-wafer-scale-engine-why-we-need-big-chips-for-deep-learning/>

Cerebras's Wafer Scale Engine-2 (2021)



Cerebras WSE-2
2.6 Trillion transistors
46,225 mm²

- The largest ML accelerator chip
- 850,000 cores




Largest GPU
54.2 Billion transistors
826 mm²

NVIDIA Ampere GA100

<https://www.anandtech.com/show/14758/hot-chips-31-live-blogs-cerebras-wafer-scale-deep-learning>

<https://www.cerebras.net/cerebras-wafer-scale-engine-why-we-need-big-chips-for-deep-learning>

Cerebras's Wafer Scale Engine-3 (2023)



Cerebras Wafer-Scale Engine

The largest chip ever produced

46,225 mm² silicon

4 trillion transistors

900,000 AI cores

125 Petaflops of AI compute

44 Gigabytes of on-chip memory

21 PByte/s memory bandwidth

214 Pbit/s fabric bandwidth


5nm TSMC process

Huge Demand for Performance & Efficiency

Exponential Growth of Neural Networks



Memory and compute requirements




1800x more compute
In just **2 years**

Tomorrow, **multi-trillion**
parameter models

More on the Cerebras WSE

<https://www.youtube.com/watch?v=x2-qB0J7KHw>



The image shows a YouTube video player interface. The main video frame displays a title card with a background of a circuit board. The title is "Thinking Outside the Die: Architecting the ML Accelerator of the Future" in large white text. Below the title, the speaker's name "Sean Lie" and his role "Co-founder & Chief HW Architect, Cerebras" are listed in white text. In the top right corner of the video frame, there is a small inset portrait of Sean Lie, a man with glasses wearing a light green shirt. At the bottom left of the video frame, there is a black bar with white text that says "Live in 9 days February 28, 6:00 PM" and a bell icon with the text "Reminder on". Below the video frame, the video title "SAFARI Live Seminar - Thinking Outside the Die: Architecting the ML Accelerator of the Future" is displayed. Below the title, it says "1 waiting • Scheduled for Feb 28, 2022". To the right of the title, there are icons for likes (7), dislikes, shares, and a save button. At the bottom left, there is a channel logo for "SAI" and a profile picture of Onur Mutlu, with the text "Onur Mutlu Lectures" and "22.6K subscribers". At the bottom right, there are two blue buttons labeled "ANALYTICS" and "EDIT VIDEO".

Thinking Outside the Die:
Architecting the ML Accelerator of the Future

Sean Lie
Co-founder & Chief HW Architect, Cerebras

Live in 9 days
February 28, 6:00 PM

Reminder on

SAFARI Live Seminar - Thinking Outside the Die: Architecting the ML Accelerator of the Future

1 waiting • Scheduled for Feb 28, 2022

7 DISLIKE SHARE SAVE ...

SAI Onur Mutlu Lectures
22.6K subscribers

ANALYTICS EDIT VIDEO

Microsoft Brainwave FPGA Acceleration

Serving DNNs in Real Time at Datacenter Scale with Project Brainwave

Eric Chung, Jeremy Fowers, Kalin Ovtcharov, Michael Papamichael, Adrian Caulfield, Todd Massengill, Ming Liu, Daniel Lo, Shlomi Alkalay, Michael Haselman, Maleen Abeysdeera, Logan Adams, Hari Angepat, Christian Boehn, Derek Chiou, Oren Firestein, Alessandro Forin, Kang Su Gatlin, Mahdi Ghandi, Stephen Heil, Kyle Holohan, Ahmad El Husseini, Tamas Juhasz, Kara Kagi, Ratna K. Kovvuri, Sitaram Lanka, Friedel van Megan, Dima Mukhortov, Prerak Patel, Brandon Perez, Amanda Grace Rapsang, Steven K. Reinhardt, Bitu Darvish Rouhani, Adam Sapek, Raja Seera, Sangeetha Shekar, Balaji Sridharan, Gabriel Weisz, Lisa Woods, Phillip Yi Xiao, Dan Zhang, Ritchie Zhao, and Doug Burger
Microsoft Corporation

To meet the computational demands required of deep learning, cloud operators are turning toward specialized hardware for improved efficiency and performance. Project Brainwave, Microsoft's principal infrastructure for AI serving in real time, accelerates deep neural network (DNN) inferencing in major services such as Bing's intelligent search features and Azure. Exploiting distributed model parallelism and pinning over low-latency hardware microservices, Project Brainwave serves state-of-the-art, pre-trained DNN models with high efficiencies at low batch sizes. A high-performance, precision-adaptable FPGA soft processor is at the heart of the system, achieving up to 39.5 teraflops (Tflops) of effective performance at Batch 1 on a state-of-the-art Intel Stratix 10 FPGA.

The recent successes of machine learning, enabled largely by the rise of DNNs, have fueled a growing demand for ubiquitous AI, ranging from conversational agents to object recognition to intelligent search. While state-of-the-art DNNs continue to deliver major breakthroughs in challenging AI domains such as computer vision and natural language processing, their computational demands have

steadily outpaced the performance growth rate of standard CPUs. These trends have spurred a




Figure 3. The framework-neutral Brainwave tool flow accepts models from different DNN toolchains and exports them into a common intermediate graph representation. Tool flow optimizes the intermediate representation and partitions it into sub-graphs assigned to different CPUs and FPGAs. Device-specific backends generate device assembly and are linked together by a federated runtime that gets deployed into a live FPGA hardware microservice.




Figure 5. An independent SRAM memory port is dedicated to every lane of a multi-lane vector dot product unit within the MVU, allowing up to 80,000 MACs on a Stratix 10 280 to be fed with independent weights. As a result, FPGA can achieve near-peak utilization on Batch 1-oriented matrix-vector multiplication.

Food for Thought

- Compare and contrast **GPUs** vs **Systolic Arrays** (vs **FPGAs**)
 - Which one is better for machine learning?
 - Which one is better for image/vision processing?
 - What types of parallelism does each one exploit?
 - What are the tradeoffs?
- If you are interested in such questions and more, take:
 - **Seminar in Computer Architecture**
 - **Computer Architecture Master's Course**

Digital Design & Computer Arch.

Lecture 18b: Systolic Array Architectures

Prof. Onur Mutlu

ETH Zürich

Spring 2026


24 April 2026

Backup Slides

(for Further Study)


Google Edge Neural Network Models

We analyze inference execution using 24 edge NN models



Diversity Across the Models



Insight 1: there is **significant variation** in terms of layer characteristics **across the models**



Diversity Within the Models

Insight 2: even **within** each model, layers exhibit **significant variation** in terms of layer characteristics

For example, our analysis of edge **CNN** models shows:




Variation in **MAC intensity**: up to **200x** across layers

Variation in **FLOP/Byte**: up to **244x** across layers


Mensa High-Level Overview

Edge TPU Accelerator



Model A



Model B




Model C




Monolithic Accelerator

Mensa


Model A



Model B




Model C






Family 1




Family 2




Family 3



Acc. 1



Acc. 2




Acc. 3

Heterogeneous Accelerators

Identifying Layer Families


Key observation: the majority of layers group into a small number of layer families



Families 1 & 2: low parameter footprint, high data reuse and **MAC intensity**
→ compute-centric layers


Families 3, 4 & 5: high parameter footprint, low data reuse and **MAC intensity**
→ data-centric layers

Mensa: Energy Reduction



Mensa-G reduces energy consumption by 3.0X
compared to the baseline Edge TPU

Mensa: Throughput Improvement



Mensa-G improves inference throughput by 3.1X
compared to the baseline Edge TPU

Mensa: Highly-Efficient ML Inference

- Amirali Boroumand, Saugata Ghose, Berkin Akin, Ravi Narayanaswami, Geraldo F. Oliveira, Xiaoyu Ma, Eric Shiu, and Onur Mutlu,
"Google Neural Network Models for Edge Devices: Analyzing and Mitigating Machine Learning Inference Bottlenecks"
Proceedings of the 30th International Conference on Parallel Architectures and Compilation Techniques (PACT), Virtual, September 2021.
[[Slides \(pptx\)](#)] [[pdf](#)]
[[Talk Video](#) (14 minutes)]

Google Neural Network Models for Edge Devices: Analyzing and Mitigating Machine Learning Inference Bottlenecks

Amirali Boroumand^{†◇}

Geraldo F. Oliveira^{*}

Saugata Ghose[‡]

Xiaoyu Ma[§]

Berkin Akin[§]

Eric Shiu[§]

Ravi Narayanaswami[§]

Onur Mutlu^{*†}

[†]*Carnegie Mellon Univ.*

[◇]*Stanford Univ.*


[‡]*Univ. of Illinois Urbana-Champaign*

[§]*Google*

^{*}*ETH Zürich*

A MIMD Machine with SIMD Processors (I)

- **MIMD** machine
 - ❑ Distributed memory (no shared memory)
 - ❑ 2D-mesh interconnection fabric



A MIMD Machine with SIMD Processors (II)

- SIMD processors
 - ❑ 4-way SIMD for 16-bit floating point operands
 - ❑ 48 KB of local SRAM

